

Part I - (Loan Data from Prosper)

by Chia-Hung Lee

Introduction

This data set contains 113,937 loans with 81 variables on each loan, including loan amount, borrower rate (or interest rate), current loan status, borrower income, and many others.

Preliminary Wrangling

```
In [1]: # import all packages and set plots to be embedded inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

```
In [2]: df = pd.read_csv('Loan-Data-from-Prosper.csv')
print('Dimensions of the dataset = {}'.format(df.shape))
print('-----')
print(df.dtypes)
```

Dimensions of the dataset = (113937, 81)

```
-----
ListingKey                object
ListingNumber             int64
ListingCreationDate       object
CreditGrade              object
Term                     int64
...
PercentFunded            float64
Recommendations          int64
InvestmentFromFriendsCount int64
InvestmentFromFriendsAmount float64
Investors                int64
Length: 81, dtype: object
```

```
In [3]: df.head()
```

Out [3]:

	ListingKey	ListingNumber	ListingCreationDate	CreditGrade	Term	LoanSta
0	1021339766868145413AB3B	193129	2007-08-26 19:09:29.263000000	C	36	Comple
1	10273602499503308B223C1	1209647	2014-02-27 08:28:07.900000000	NaN	36	Curi
2	0EE9337825851032864889A	81716	2007-01-05 15:00:47.090000000	HR	36	Comple
3	0EF5356002482715299901A	658116	2012-10-22 11:02:35.010000000	NaN	36	Curi
4	0F023589499656230C5E3E2	909464	2013-09-14 18:38:39.097000000	NaN	36	Curi

5 rows x 81 columns

In [4]:

```
print(df.dtypes.unique())
print('-----')
print('Object features: {}'.format(df.dtypes[df.dtypes=='0'].count()))
print(df.dtypes[df.dtypes=='0'])
print('-----')
print('Object features: {}'.format(df.dtypes[df.dtypes=='int64'].count()))
print(df.dtypes[df.dtypes=='int64'])
print('-----')
print('Object features: {}'.format(df.dtypes[df.dtypes=='float64'].count()))
print(df.dtypes[df.dtypes=='float64'])
print('-----')
print('Object features: {}'.format(df.dtypes[df.dtypes=='bool'].count()))
print(df.dtypes[df.dtypes=='bool'])
```

```
[dtype('O') dtype('int64') dtype('float64') dtype('bool')]
```

Object features: 17

ListingKey	object
ListingCreationDate	object
CreditGrade	object
LoanStatus	object
ClosedDate	object
ProsperRating (Alpha)	object
BorrowerState	object
Occupation	object
EmploymentStatus	object
GroupKey	object
DateCreditPulled	object
FirstRecordedCreditLine	object
IncomeRange	object
LoanKey	object
LoanOriginationDate	object
LoanOriginationQuarter	object
MemberKey	object

dtype: object

Object features: 11

ListingNumber	int64
Term	int64
ListingCategory (numeric)	int64
OpenRevolvingAccounts	int64
LoanCurrentDaysDelinquent	int64
LoanMonthsSinceOrigination	int64
LoanNumber	int64
LoanOriginalAmount	int64
Recommendations	int64
InvestmentFromFriendsCount	int64
Investors	int64

dtype: object

Object features: 50

BorrowerAPR	float64
BorrowerRate	float64
LenderYield	float64
EstimatedEffectiveYield	float64
EstimatedLoss	float64
EstimatedReturn	float64
ProsperRating (numeric)	float64
ProsperScore	float64
EmploymentStatusDuration	float64
CreditScoreRangeLower	float64
CreditScoreRangeUpper	float64
CurrentCreditLines	float64
OpenCreditLines	float64
TotalCreditLinespast7years	float64
OpenRevolvingMonthlyPayment	float64
InquiriesLast6Months	float64
TotalInquiries	float64
CurrentDelinquencies	float64
AmountDelinquent	float64
DelinquenciesLast7Years	float64
PublicRecordsLast10Years	float64
PublicRecordsLast12Months	float64
RevolvingCreditBalance	float64
BankcardUtilization	float64
AvailableBankcardCredit	float64
TotalTrades	float64
TradesNeverDelinquent (percentage)	float64
TradesOpenedLast6Months	float64

DebtToIncomeRatio	float64
StatedMonthlyIncome	float64
TotalProsperLoans	float64
TotalProsperPaymentsBilled	float64
OnTimeProsperPayments	float64
ProsperPaymentsLessThanOneMonthLate	float64
ProsperPaymentsOneMonthPlusLate	float64
ProsperPrincipalBorrowed	float64
ProsperPrincipalOutstanding	float64
ScoreExchangeAtTimeOfListing	float64
LoanFirstDefaultedCycleNumber	float64
MonthlyLoanPayment	float64
LP_CustomerPayments	float64
LP_CustomerPrincipalPayments	float64
LP_InterestandFees	float64
LP_ServiceFees	float64
LP_CollectionFees	float64
LP_GrossPrincipalLoss	float64
LP_NetPrincipalLoss	float64
LP_NonPrincipalRecoverypayments	float64
PercentFunded	float64
InvestmentFromFriendsAmount	float64
dtype:	object

Object features: 3

IsBorrowerHomeowner	bool
CurrentlyInGroup	bool
IncomeVerifiable	bool

dtype: object

Note on data description

- ListingKey, ListingNumber: are ID for listing, could be used as index.
- CreditGrade: Only for listings in and before 2009.
- LoanStatus: [Cancelled, Chargedoff, Completed, Current, Defaulted, FinalPaymentInProgress, PastDue.]
- [EstimatedEffectiveYield, EstimatedLoss, EstimatedReturn]: Only for listings after July 2009.
- [ProsperRating(numeric), ProsperRating(Alpha), ProsperScore]: Only for listings after July 2009.
- [CurrentlyInGroup, GroupKey]: affiliation to specific groups.

What is the structure of your dataset?

There are 113,937 load data in the dataset with 81 features (Object: 17; Int64: 11; Float64: 50; Boolean:3). The majority (61 variables) are numeric in nature.

What is/are the main feature(s) of interest in your dataset?

While the dataset offers an array of features for exploration, this analysis focuses primarily on investigating the BorrowerAPR and BorrowerRate variables, in addition to other relevant attributes. We are trying to answer the following questions:

- What affects the borrower's APR or interest rate?
- What affects the original loan amount?

What features in the dataset do you think will help support your investigation into your feature(s) of interest?

The key attributes that will play a pivotal role in supporting the analysis of BorrowerAPR and BorrowerRate include:

Column	Description
--------	-------------

|Term | The length of the loan expressed in months. **|LoanStatus** | The current status of the loan: Cancelled, Chargedoff, Completed, Current, Defaulted, FinalPaymentInProgress, PastDue. The PastDue status will be accompanied by a delinquency bucket. **|BorrowerAPR** | The Borrower's Annual Percentage Rate (APR) for the loan. **|BorrowerRate** | The Borrower's interest rate for this loan. **|ProsperRating (Alpha)** | The Prosper Rating assigned at the time the listing was created between AA - HR. Applicable for loans originated after July 2009. **|ListingCategory** | The category of the listing that the borrower selected when posting their listing: 0 - Not Available, 1 - Debt Consolidation, 2 - Home Improvement, 3 - Business, 4 - Personal Loan, 5 - Student Use, 6 - Auto, 7- Other, 8 - Baby&Adoption, 9 - Boat, 10 - Cosmetic Procedure, 11 - Engagement Ring, 12 - Green Loans, 13 - Household Expenses, 14 - Large Purchases, 15 - Medical/Dental, 16 - Motorcycle, 17 - RV, 18 - Taxes, 19 - Vacation, 20 - Wedding Loans **|EmploymentStatus** | The employment status of the borrower at the time they posted the listing. **|IsBorrowerHomeowner** | A Borrower will be classified as a homeowner if they have a mortgage on their credit profile or provide documentation confirming they are a homeowner. TRUE or FALSE **|IncomeRange** | The income range of the borrower at the time the listing was created. **|DebtToIncomeRatio** | The debt to income ratio of the borrower at the time the credit profile was pulled. This value is Null if the debt to income ratio is not available. This value is capped at 10.01 (any debt to income ratio larger than 1000% will be returned as 1001%). **|StatedMonthlyIncome** | The monthly income the borrower stated at the time the listing was created. **|LoanOriginalAmount** | The original amount of the loan. **|MonthlyLoanPayment** | The scheduled monthly loan payment.

Constructing the Modified Dataset to for Analysis

```
In [5]: df_mod = df[['LoanKey','Term','LoanStatus','BorrowerAPR','BorrowerRate',
                    'ProsperRating (Alpha)','ListingCategory (numeric)','EmploymentStatus',
                    'IsBorrowerHomeowner','IncomeRange','DebtToIncomeRatio',
                    'LoanOriginalAmount','StatedMonthlyIncome','MonthlyLoanPayment']].copy()
df_mod.head()
```

```
Out [5]:
```

	LoanKey	Term	LoanStatus	BorrowerAPR	BorrowerRate	ProsperRating (Alpha)
0	E33A3400205839220442E84	36	Completed	0.16516	0.1580	NaN
1	9E3B37071505919926B1D82	36	Current	0.12016	0.0920	A
2	6954337960046817851BCB2	36	Completed	0.28269	0.2750	NaN
3	A0393664465886295619C51	36	Current	0.12528	0.0974	A
4	A180369302188889200689E	36	Current	0.24614	0.2085	D

```
In [6]: # Rename columns
df_mod.rename(columns={'ProsperRating (Alpha)': 'ProsperRating', 'ListingCategory (num
df_mod.head()
```

Out [6]:	LoanKey	Term	LoanStatus	BorrowerAPR	BorrowerRate	ProsperRating
0	E33A3400205839220442E84	36	Completed	0.16516	0.1580	NaN
1	9E3B37071505919926B1D82	36	Current	0.12016	0.0920	A
2	6954337960046817851BCB2	36	Completed	0.28269	0.2750	NaN
3	A0393664465886295619C51	36	Current	0.12528	0.0974	A
4	A180369302188889200689E	36	Current	0.24614	0.2085	D

```
In [7]: df_mod.duplicated().value_counts()
```

```
Out[7]: False      113066
        True        871
        dtype: int64
```

```
In [8]: # Remove duplicate data
df_mod = df_mod.drop_duplicates()
df_mod.duplicated().value_counts()
```

```
Out[8]: False      113066
        dtype: int64
```

```
In [9]: df_mod.isnull().sum()
```

```
Out[9]: LoanKey          0
        Term            0
        LoanStatus      0
        BorrowerAPR     25
        BorrowerRate    0
        ProsperRating   29084
        ListingCategory  0
        EmploymentStatus 2255
        IsBorrowerHomeowner 0
        IncomeRange     0
        DebtToIncomeRatio 8472
        LoanOriginalAmount 0
        StatedMonthlyIncome 0
        MonthlyLoanPayment 0
        dtype: int64
```

```
In [10]: df_mod=df_mod.dropna()
df_mod.isnull().sum()
```

```
Out[10]: LoanKey          0
        Term            0
        LoanStatus      0
        BorrowerAPR     0
        BorrowerRate    0
        ProsperRating   0
        ListingCategory  0
        EmploymentStatus 0
        IsBorrowerHomeowner 0
        IncomeRange     0
        DebtToIncomeRatio 0
        LoanOriginalAmount 0
        StatedMonthlyIncome 0
        MonthlyLoanPayment 0
        dtype: int64
```

```
In [11]: df_mod.set_index('LoanKey',inplace=True)
```

```
In [12]: df_mod.shape
```

```
Out[12]: (76768, 13)
```

Structure of my dataset (df_mod)

Number of Rows in dataset are 76768 and Number of Columns in dataset are 13.

What is/are the main feature(s) of interest in your dataset?

- LoanKey (Index)
- Term
- LoanStatus
- BorrowerAPR
- BorrowerRate
- ProsperRating (Alpha)
- ListingCategory (numeric)
- EmploymentStatus
- IsBorrowerHomeowner
- IncomeRange
- DebtToIncomeRatio
- LoanOriginalAmount
- StatedMonthlyIncome
- MonthlyLoanPayment

```
In [13]: df_mod.head()
```

	Term	LoanStatus	BorrowerAPR	BorrowerRate	ProsperRating	Li
LoanKey						
9E3B37071505919926B1D82	36	Current	0.12016	0.0920		A
A0393664465886295619C51	36	Current	0.12528	0.0974		A
A180369302188889200689E	36	Current	0.24614	0.2085		D
C3D63702273952547E79520	60	Current	0.15425	0.1314		B
CE963680102927767790520	36	Current	0.31032	0.2712		E

Univariate Exploration

Term

```
In [14]: df_mod.Term.value_counts()
```

```
Out[14]: 36    52505
        60    22849
        12     1414
        Name: Term, dtype: int64
```

```
In [15]: # Plot a countplot for the distribution of loan terms
plt.figure(figsize=[8,5])
```


```
ax = sns.countplot(data=df_mod, x='Term', palette = 'Blues', edgecolor='black')

plt.xlabel('Loan Term (Months)') # Label for the x-axis
plt.ylabel('Count') # Label for the y-axis
plt.title('Distribution of Loan Terms') # Title for the plot

# Show percentage labels
total_count = len(df_mod)
for i in ax.patches:
    percentage = '{:.1f}%'.format(100 * i.get_height() / total_count)
    x = i.get_x() + i.get_width() / 2 - 0.1
    y = i.get_height() + 20
    ax.annotate(percentage, (x, y), fontsize=10, ha='center')

# Rotate x-axis labels for better readability
plt.xticks(rotation=45)

plt.show() # Display the plot
```



- There are **3 available loan term**. The most frequently selected term is **36 month**, although there is also a notable preference for the **60 month** term among some borrowers.

LoanStatus

- Which categories do majority of loans fall in ?

```
In [16]: df_mod['LoanStatus'].value_counts()
```



```
Out [16]: Current          51712
Completed        17691
Chargedoff       4445
Defaulted        885
Past Due (1-15 days)  716
Past Due (31-60 days)  325
Past Due (91-120 days) 277
Past Due (61-90 days) 274
Past Due (16-30 days) 242
FinalPaymentInProgress 187
Past Due (>120 days)  14
Name: LoanStatus, dtype: int64
```

```
In [17]: def BarPlot(df, var, order=None, figsize=[12,6], title=''):
# Create the figure and axes
fig, ax = plt.subplots(figsize=figsize)

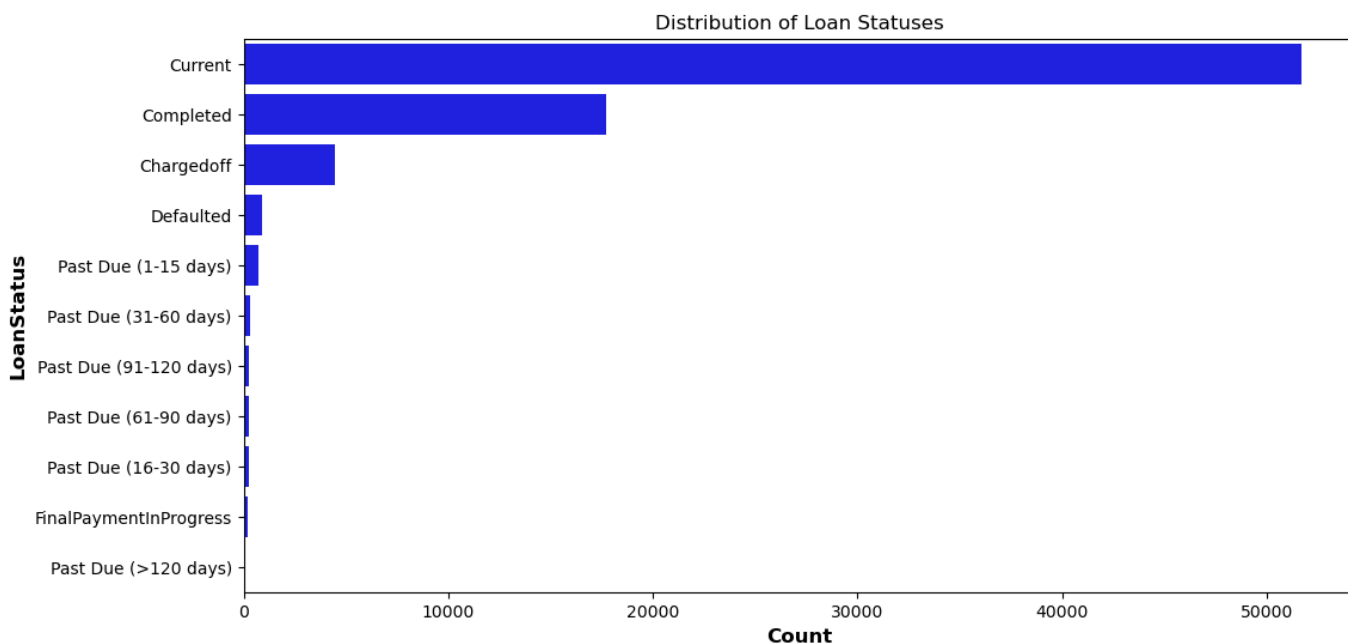
# Plot the barplot
type_count = df[var].value_counts()
sns.barplot(x=type_count, y=type_count.index, order=order, color='blue')

# Add title
ax.set_title(title)

# Add X and Y labels and format them
ax.set_xlabel('Count', fontsize=12, weight="bold")
ax.set_ylabel(var, fontsize=12, weight="bold")

plt.show()
```

```
In [18]: BarPlot(df_mod, 'LoanStatus', title='Distribution of Loan Statuses')
```



- The plot above illustrates that the majority of loans within the dataset belong to the categories of "Current," followed by "Completed," "Charged-off," "Defaulted," and "Past Due (1-15 days)."

BorrowerAPR

```
In [19]: df_mod.BorrowerAPR.describe()
```

```
Out [19]: count      76768.000000
mean         0.223978
std          0.079291
min          0.045830
25%          0.162590
50%          0.215660
75%          0.287800
max          0.423950
Name: BorrowerAPR, dtype: float64
```

```
In [20]: # Function to plot Histogram distribution
def HistPlot(df, var, interval=30, figsize=[12,6], title=''):
    # Set intervals for bins
    bins = np.arange(df[var].min(), df[var].max() + interval, interval)

    # Create the figure and axes
    fig, ax = plt.subplots(figsize=figsize)

    # Plot the histogram
    ax.hist(df[var], bins=bins, edgecolor='black', color='Blue', alpha=0.7)

    # Add a vertical line for the median
    median = df[var].median()
    ax.axvline(median, color='red', linestyle='dashed', linewidth=2, label=f'Median:

    # Add title
    ax.set_title(title)


    # Add X and Y labels and format them
    ax.set_xlabel(var, fontsize=12, weight="bold")
    ax.set_ylabel('Frequency', fontsize=12, weight="bold")

    # Add grid lines for better readability
    ax.grid(axis='y', linestyle='--', alpha=0.7)

    # Add a legend for the median line
    ax.legend(loc='upper right')

    plt.show()
```

```
In [21]: # Call plot function (using all data)
HistPlot(df_mod, 'BorrowerAPR', 0.01)
```



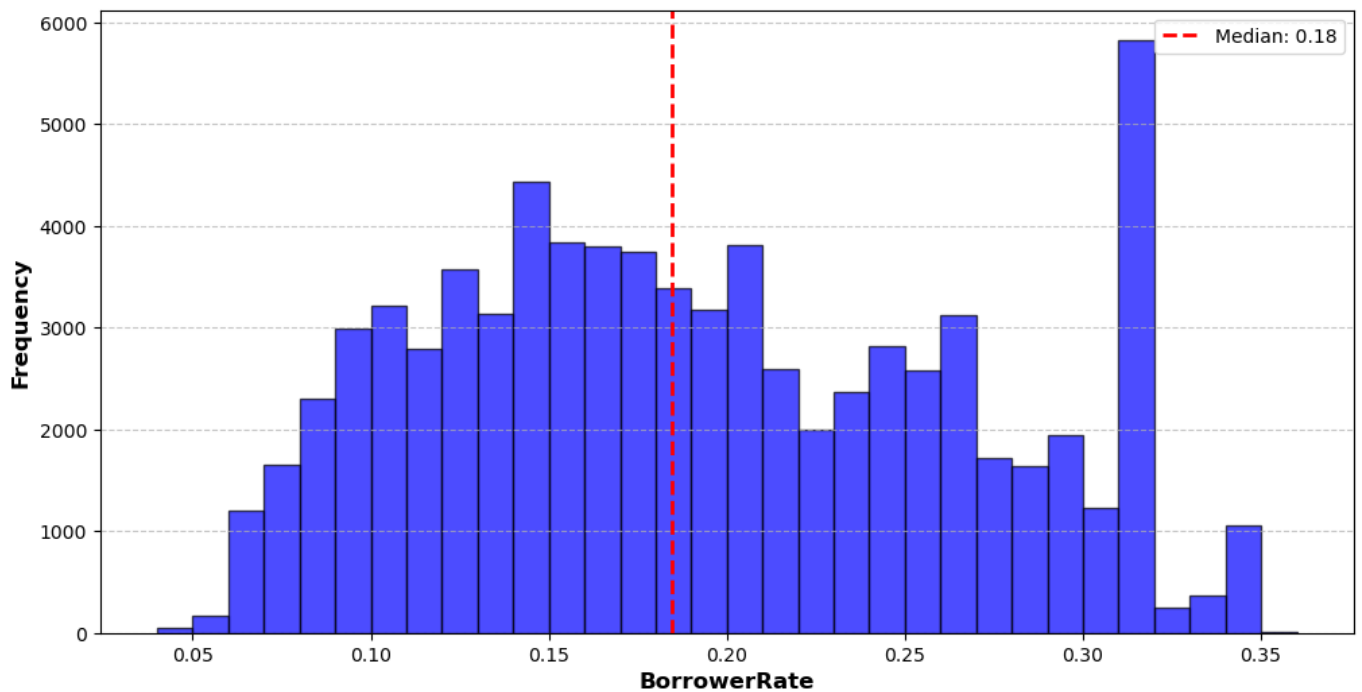
- **Higher BorrowerAPR** indicates a greater interest rate on borrowed funds, resulting in increased interest payments.
- In the distribution plot of BorrowerAPR, there is a prominent peak around 0.09, 0.23, 0.25, 0.27, 0.29, a minor peak at approximately 0.18, and a significant peak around 0.36. Very few individuals have an APR exceeding 0.4.

BorrowerRate

```
In [22]: df_mod.BorrowerRate.describe()
```

```
Out[22]: count      76768.000000
mean         0.193653
std          0.074018
min          0.040000
25%          0.134900
50%          0.184500
75%          0.254900
max          0.360000
Name: BorrowerRate, dtype: float64
```

```
In [23]: HistPlot(df_mod, 'BorrowerRate', 0.01)
```



- The distribution of **BorrowerRate** is multimodal, with a prominent peak at approximately 0.32. Rates exceeding 0.35 are rare, with a median rate of 0.18.

ProsperRating

```
In [24]: df_mod.ProsperRating.value_counts()
```

```
Out[24]: C      16671
B      14444
A      13555
D      12724
E       8543
HR       5722
AA       5109
Name: ProsperRating, dtype: int64
```



```

9 : 'Boat',
10 : 'Cosmetic Procedure',
11 : 'Engagement Ring',
12 : 'Green Loans',
13 : 'Household Expenses',
14 : 'Large Purchases',
15 : 'Medical/Dental',
16 : 'Motorcycle',
17 : 'RV',
18 : 'Taxes',
19 : 'Vacation',
20 : 'Wedding Loans'}

```

```
In [28]: df_mod['ListingCategory'] = df_mod['ListingCategory'].replace(to_replace=listing_cate
```

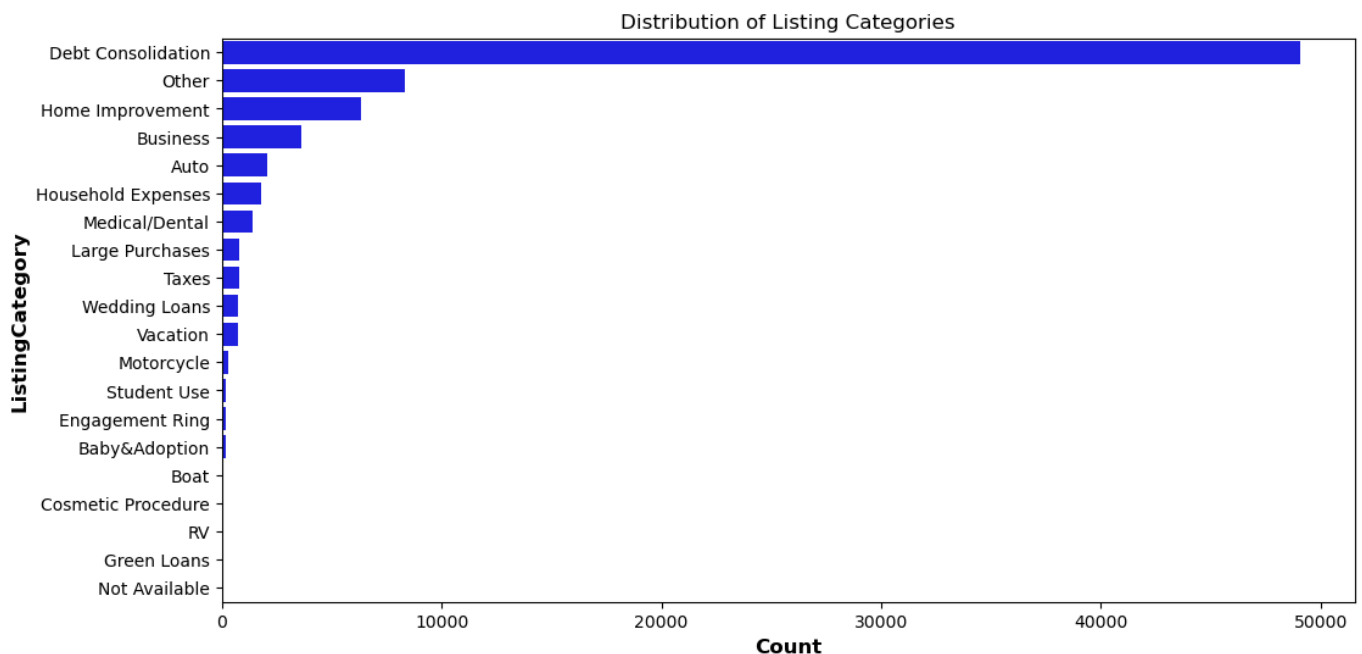
```
In [29]: df_mod.ListingCategory.value_counts()
```

```

Out[29]: Debt Consolidation    49099
Other                        8334
Home Improvement            6326
Business                   3626
Auto                       2038
Household Expenses         1779
Medical/Dental             1390
Large Purchases             794
Taxes                      785
Wedding Loans              724
Vacation                   718
Motorcycle                 289
Student Use                201
Engagement Ring            198
Baby&Adoption              188
Boat                       83
Cosmetic Procedure         82
RV                         50
Green Loans                45
Not Available               19
Name: ListingCategory, dtype: int64

```

```
In [30]: BarPlot(df_mod, 'ListingCategory', title='Distribution of Listing Categories')
```



- From the chart above, it is evident that a significant number of individuals acquire loans for the purpose of **"Debt Consolidation"** primarily to manage and

EmploymentStatus

```
In [31]: df_mod.EmploymentStatus.value_counts()
```

```
Out[31]: Employed          65160
Full-time          7584
Other             3462
Retired           320
Part-time         199
Self-employed      42
Not employed       1
Name: EmploymentStatus, dtype: int64
```

```
In [32]: # Calculate the percentage of borrowers in each employment status category
percentage_data = (df_mod['EmploymentStatus'].value_counts() / len(df_mod)) * 100

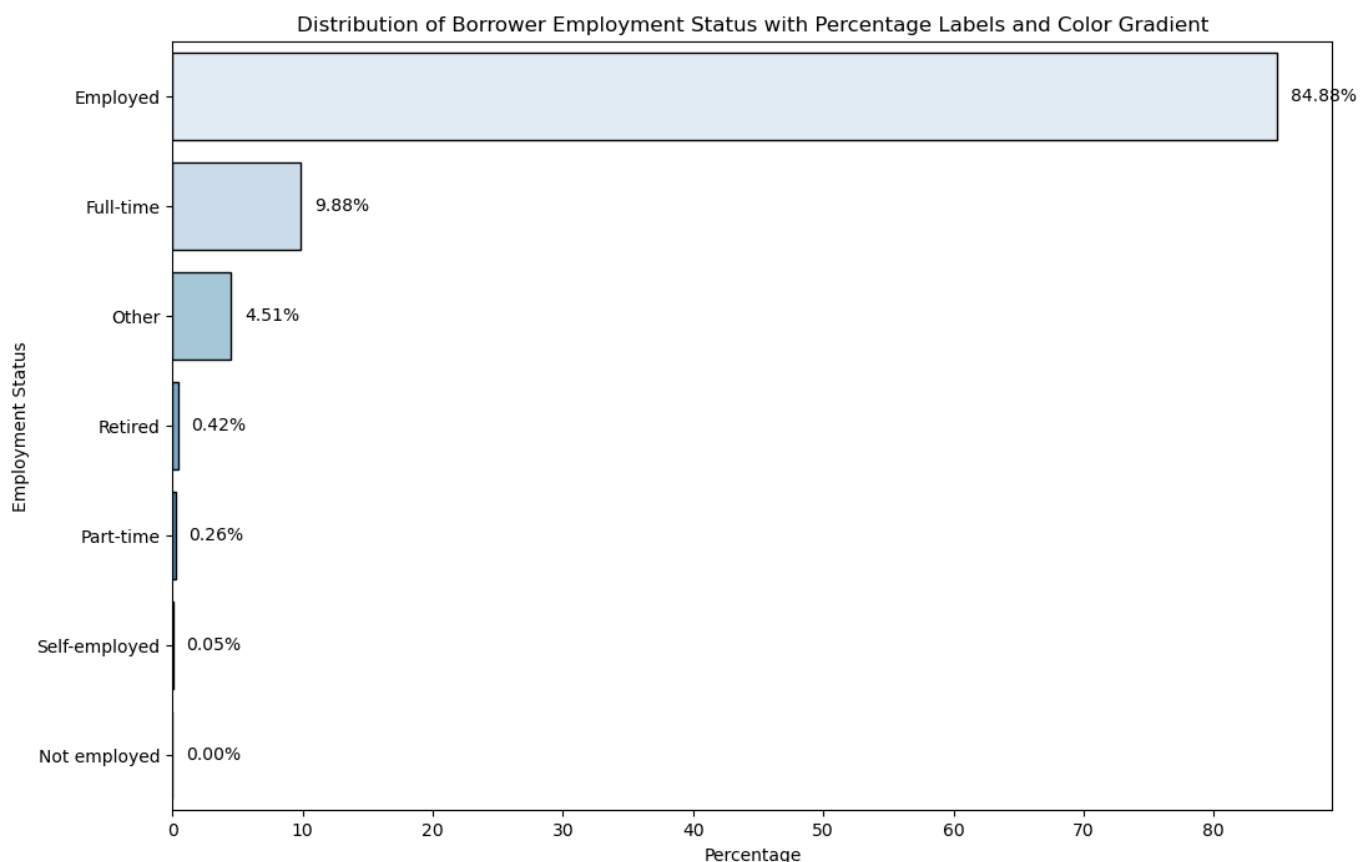
# Create a color palette with a gradient
color_palette = sns.color_palette("Blues", len(percentages_data))

# Sort the data by count to maintain color consistency
percentage_data = percentage_data.sort_values(ascending=False)

# Create the bar chart with color gradient
plt.figure(figsize=(12, 8))
bars = sns.barplot(y=percentage_data.index, x=percentage_data.values, palette=color_p

# Add percentage labels on the bars
for bar, percentage in zip(bars.patches, percentage_data):
    width = bar.get_width()
    plt.text(width + 1, bar.get_y() + bar.get_height() / 2, f'{percentage:.2f}%', ha=

plt.xlabel('Percentage')
plt.ylabel('Employment Status')
plt.title('Distribution of Borrower Employment Status with Percentage Labels and Colo
plt.show()
```




- The majority of loan borrowers are **Employed**, which is a logical trend as a source of income is necessary to repay the loans.

IsBorrowerHomeowner

```
In [33]: df_mod.IsBorrowerHomeowner.value_counts()
```

```
Out[33]: True      41676  
        False    35092  
        Name: IsBorrowerHomeowner, dtype: int64
```

```
In [34]: # Pie chart to show the proportion of people who are homeowners  
labels = ['Yes', 'No'] # Define the labels  
plt.figure(figsize=(8, 8))  
colors = ['lightgreen', 'lightcoral']  
explode = (0.1, 0) # Explode the 'Yes' slice for emphasis  
  
plt.pie(df_mod['IsBorrowerHomeowner'].value_counts(), labels=labels, autopct='%0.1f%%',  
        colors=colors, explode=explode)  
plt.title('Proportion of Homeowners among Borrowers')  
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.  
  
plt.show()
```



- A slight majority of borrowers are homeowners.

IncomeRange

```
In [35]: df_mod.IncomeRange.value_counts()
```

```
Out[35]: $50,000–74,999    23756
          $25,000–49,999    21795
          $100,000+         13889
          $75,000–99,999    13519
          $1–24,999         3808
          Not employed       1
          Name: IncomeRange, dtype: int64
```

```
In [36]: # Create ordinal categories for income ranges
ordinal_income_ranges = ['$100,000+', '$75,000–99,999', '$50,000–74,999', '$25,000–49,999', '$1–24,999', 'Not employed']
income_order = pd.api.types.CategoricalDtype(ordered=True, categories=ordinal_income_ranges)
df_mod['IncomeRange'] = df_mod['IncomeRange'].astype(income_order)
df_mod['IncomeRange']
```

```
Out[36]: LoanKey
9E3B37071505919926B1D82    $50,000–74,999
A0393664465886295619C51    $25,000–49,999
A180369302188889200689E    $100,000+
C3D63702273952547E79520    $100,000+
CE963680102927767790520    $25,000–49,999
...
9BD7367919051593140DB62    $50,000–74,999
62D93634569816897D5A276    $75,000–99,999
DD1A370200396006300ACA0    $25,000–49,999
589536350469116027ED11B    $25,000–49,999
00AF3704550953269A64E40    $50,000–74,999
Name: IncomeRange, Length: 76768, dtype: category
Categories (6, object): ['$100,000+' < '$75,000–99,999' < '$50,000–74,999' < '$25,000–49,999' < '$1–24,999' < '$0']
```


```
In [37]: # Plot a horizontal bar chart for the distribution of LoanStatus
plt.figure(figsize=[10, 4])

sns.countplot(data=df_mod, y='IncomeRange', palette='Blues_r', edgecolor='black') #

plt.xlabel('Count') # Label for the x-axis
plt.ylabel('Income Range') # Label for the y-axis
plt.title('Distribution of Borrower Income Range') # Title for the plot

# Display the plot
plt.show()

# Median income
median_income = df_mod['IncomeRange']
median_income = median_income.sort_values().reset_index()['IncomeRange'][76768/2]
print('The median income category = {}'.format(median_income))
```

The median income category = \$50,000–74,999

- The **median IncomeRange** (Also the most common) is \$50,000 - \$74,999 , and follows by \$25,000 - \$49,999 .

DebtToIncomeRatio

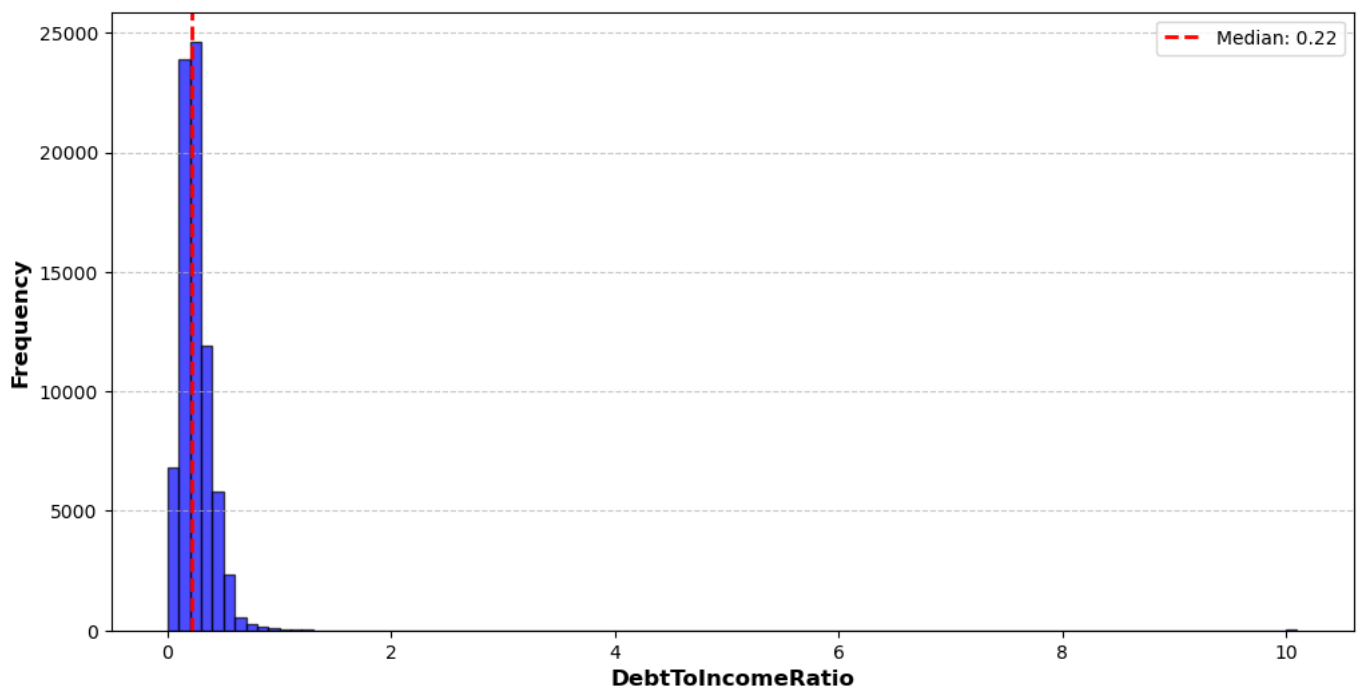
```
In [38]: df_mod.DebtToIncomeRatio.value_counts()
```

```
Out[38]: 0.18    3184
         0.22    2974
         0.17    2728
         0.14    2681
         0.21    2545
         ...
         1.54         1
         3.81         1
         4.54         1
         1.93         1
         2.53         1
         Name: DebtToIncomeRatio, Length: 259, dtype: int64
```

```
In [39]: df_mod.DebtToIncomeRatio.describe()
```

```
Out[39]: count    76768.000000
         mean      0.258692
         std       0.319727
         min       0.000000
         25%       0.150000
         50%       0.220000
         75%       0.320000
         max       10.010000
         Name: DebtToIncomeRatio, dtype: float64
```

```
In [40]: HistPlot(df_mod, 'DebtToIncomeRatio', 0.1)
```



```
In [41]: df_mod['DebtToIncomeRatio'].quantile(0.99)
```

```
Out[41]: 0.72
```

```
In [42]: x=df_mod[df_mod.DebtToIncomeRatio==10.01]
print('There are {} data entry with DebtToIncomeRatio = 10.01'.format(x.shape[0]))
```


There are 46 data entry with DebtToIncomeRatio = 10.01

- From the plot above, the majority (99%) of the **DebtToIncomeRatio** are within 1.
- There are 46 data entry with maximun DebtToIncomeRatio = 10.01 (should not be manual error)
- We will look into the majority data of the DebtToIncomeRatio

```
In [43]: fig, ax = plt.subplots(figsize=[10,6])
ax.hist(df_mod['DebtToIncomeRatio'], range=(0,1),
        bins=30, edgecolor='black', color='Blue', alpha=0.7)
# Add a vertical line at the median (50th percentile) in red
median=df_mod['DebtToIncomeRatio'].median()
plt.axvline(df_mod['DebtToIncomeRatio'].median(), color='red', linestyle='--', label=

# Remove the legend
plt.legend().set_visible(True)

plt.show()
```



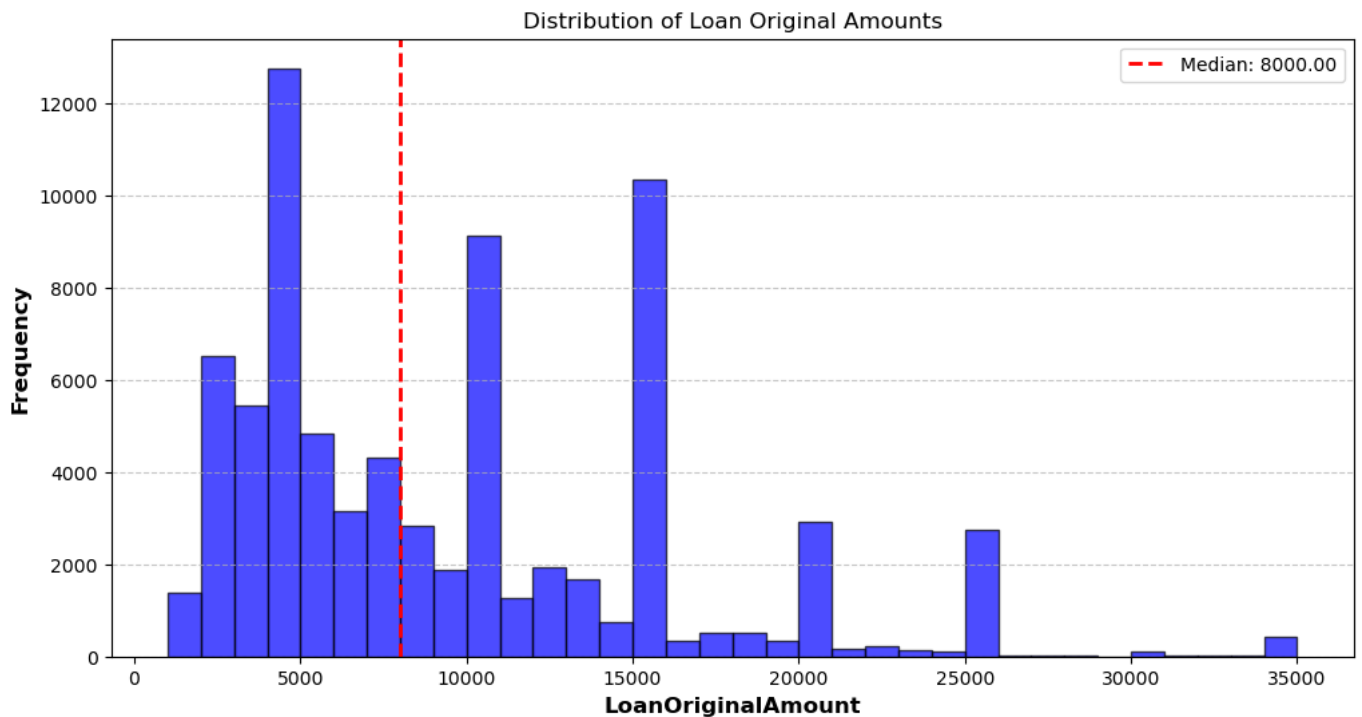
- **The majority of DebtToIncomeRatios fall below 50%**, indicating that most borrowers have a relatively low debt-to-income ratio. However, there is a notable minority with higher debt-to-income ratios. A lower ratio is generally favorable for borrowers as it signifies their ability to manage and repay their loans effectively.

OriginalLoanAmount

```
In [44]: df_mod.LoanOriginalAmount.describe()
```

```
Out [44]: count    76768.000000
mean      9248.961416
std       6389.782292
min       1000.000000
25%       4000.000000
50%       8000.000000
75%      14000.000000
max      35000.000000
Name: LoanOriginalAmount, dtype: float64
```

```
In [45]: HistPlot(df_mod, 'LoanOriginalAmount', 1000, title="Distribution of Loan Original Amo
```



- The most common **OriginalLoanAmount** occur at 4k, followed by 15k and 10k.

StatedMonthlyIncome

```
In [46]: df_mod.StatedMonthlyIncome.sort_values()
```

```
Out[46]: LoanKey
DCA9366929635721086C17C      0.250000
A2E63643277665696D47A5E      1.416667
DA623646827555431CA659F      1.833333
C9893607029050217F845B3      1.833333
B1D236532712404897E17A1      1.916667
...
1B313703263370099FC7B30     158333.333333
C6C536749708412328A7D15     394400.000000
B7F13618099831699F10189     416666.666667
5D0136156133365609F840F     466666.666667
77AC3617940949299F18FAF     483333.333333
Name: StatedMonthlyIncome, Length: 76768, dtype: float64
```

```
In [47]: df_mod.StatedMonthlyIncome.describe()
```


```
Out[47]: count      76768.000000
mean         5964.256138
std          5089.682309
min           0.250000
25%          3528.895833
50%          5000.000000
75%          7166.666667
max          483333.333333
Name: StatedMonthlyIncome, dtype: float64
```

```
In [48]: df_mod['StatedMonthlyIncome'].quantile(0.99)
```

```
Out[48]: 20416.666667
```

```
In [49]: fig, ax = plt.subplots(figsize=[10,6])
ax.hist(df_mod['StatedMonthlyIncome'],
        bins=500, edgecolor='black', color='Blue', alpha=0.7)
```

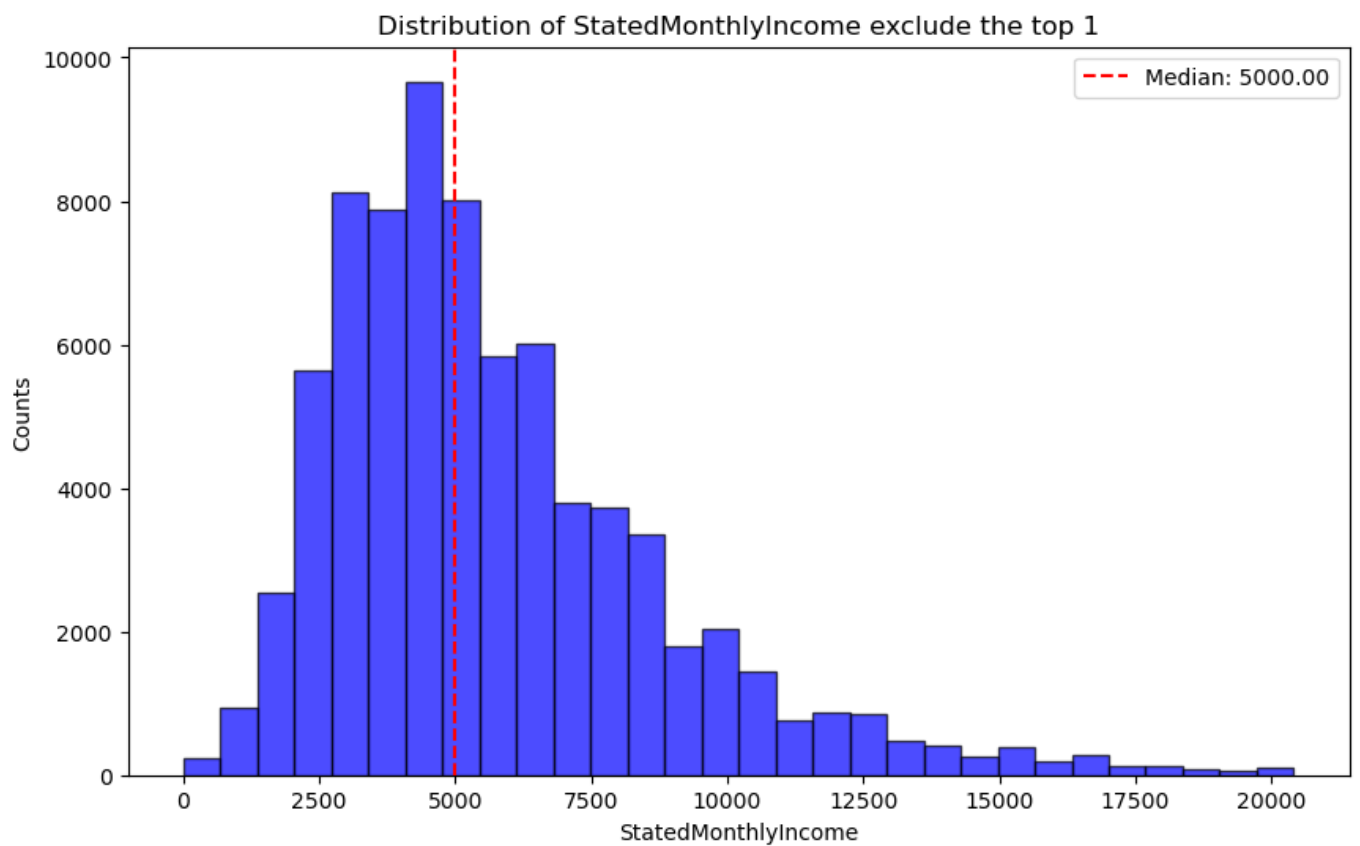
```
plt.xscale('log')
plt.show()
```



- **StatedMonthlyIncome** is highly skewed to the right.
- 99% of data are within 0 - 20526

```
In [50]: # Distribution of StatedMonthlyIncome exclude the top 1%
fig, ax = plt.subplots(figsize=[10,6])
ax.hist(df_mod['StatedMonthlyIncome'], range=(0,df_mod['StatedMonthlyIncome'].quantil
        bins=30, edgecolor='black', color='Blue', alpha=0.7)
# Add a vertical line at the median (50th percentile) in red
median=df_mod['StatedMonthlyIncome'].median()
plt.axvline(df_mod['StatedMonthlyIncome'].median(), color='red', linestyle='--', labe

# Remove the legend
plt.legend().set_visible(True)
plt.title("Distribution of StatedMonthlyIncome exclude the top 1")
plt.xlabel('StatedMonthlyIncome')
plt.ylabel('Counts')
plt.show()
```



- Most borrowers have a **StatedMonthlyIncome** that is less than 10k
- There is a peak slightly below 5k

MonthlyLoanPayment

```
In [51]: df_mod.MonthlyLoanPayment.describe()
```

```
Out[51]: count    76768.000000
mean      295.275039
std       189.109061
min        0.000000
25%       158.330000
50%       256.120000
75%       392.010000
max       2251.510000
Name: MonthlyLoanPayment, dtype: float64
```

```
In [52]: fig, ax = plt.subplots(figsize=[10,6])
ax.hist(df_mod['MonthlyLoanPayment'],
        bins=30, edgecolor='black', color='Blue', alpha=0.7)
# Add a vertical line at the median (50th percentile) in red
median=df_mod['MonthlyLoanPayment'].median()
plt.axvline(df_mod['MonthlyLoanPayment'].median(), color='red', linestyle='--', label=


# Remove the legend
plt.legend().set_visible(True)
plt.title("Distribution of MonthlyLoanPayment")
plt.xlabel('MonthlyLoanPayment')
plt.ylabel('Counts')
plt.show()
```


1. Quantitative Vs. Quantitave

```
In [54]: # Correlation plot for numeric features
plt.figure(figsize=[12, 6])

# Heatmap of the correlation matrix, annotating values.
sns.heatmap(df_mod[numeric_features].corr(), annot=True,
            fmt='.3f', center=0, cmap='Greens')

# Display the plot
plt.show()
```



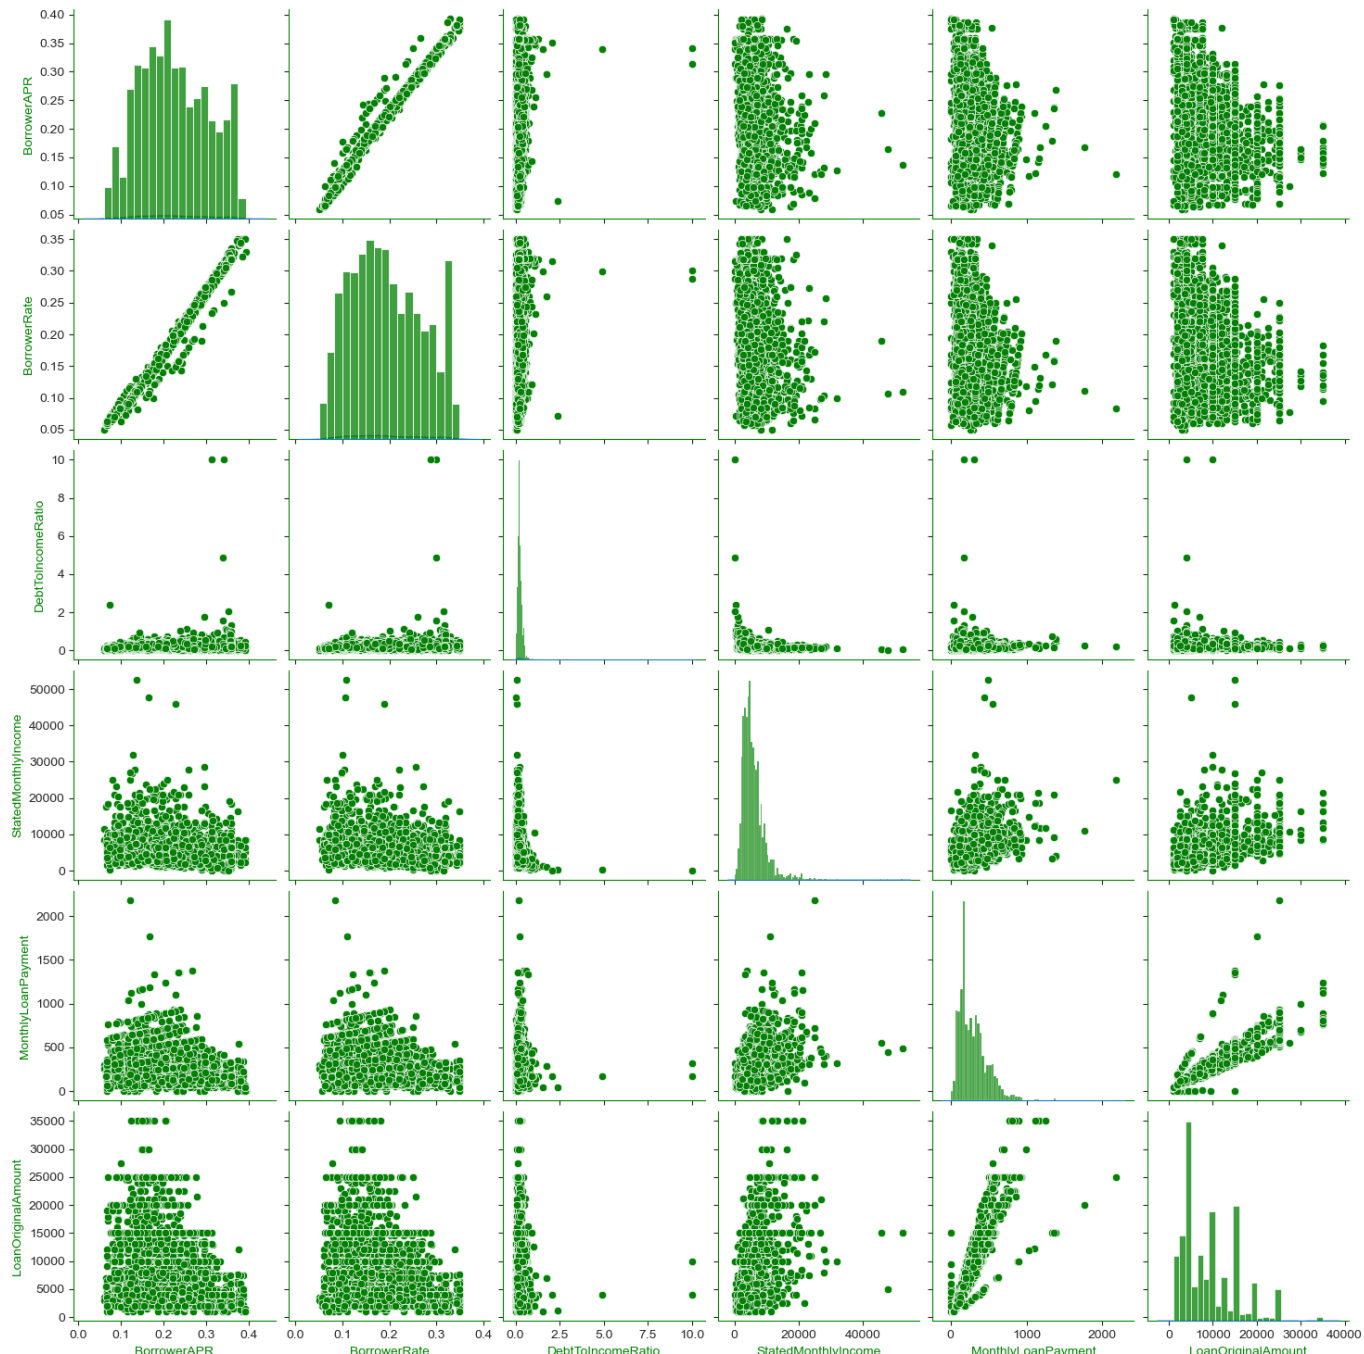
```
In [55]: # Take a sample of 3000 loans to plot
samples = np.random.choice(df_mod.index, 3000, replace=False)
loans_sample = df_mod.loc[samples, :]

# Plot pairwise relationships between all the numeric variables of interest
sns.set_style("ticks")
g = sns.pairplot(data=loans_sample, vars=numeric_features, diag_kind='kde')

# Set the color to green for all the plots
for ax in g.axes.flat:
    ax.spines['bottom'].set_color('green')
    ax.spines['top'].set_color('green')
    ax.spines['left'].set_color('green')
    ax.spines['right'].set_color('green')
    ax.xaxis.label.set_color('green')
    ax.yaxis.label.set_color('green')

# Add a grid to the plots
g.map_upper(sns.scatterplot, color='green')
g.map_lower(sns.scatterplot, color='green')
g.map_diag(sns.histplot, color='green')

plt.show()
```

- **Borrower APR** and **Borrower Rate** exhibit a robust positive correlation, with a coefficient of 0.99, which is expected since a higher APR typically leads to borrowers paying a greater amount of interest on their loans.
- **Borrower APR** and the **Loan Original Amount** demonstrate an inverse correlation with a coefficient of -0.32 . The accompanying scatter plot visually corroborates this negative association, illustrating that as the loan amount rises, the APR tends to decline.
- **Loan Original Amount** and **Stated Monthly Income** shows a slight positive correlation, as denoted by a correlation coefficient of 0.30.
- As the **Loan Original Amount** rises, there is a concurrent growth in the **Monthly Loan Payment**.


2. Quantitative vs. Qualitative

```
In [56]: # Set a larger figure size for subplots
plt.figure(figsize=[25, 10])
```

```
# Subplot 1
plt.subplot(1, 2, 1)
sns.boxplot(data=df_mod, y='IncomeRange', x='BorrowerRate', palette='Greens_r')
plt.ylabel('Income Range')
plt.xlabel('Borrower Rate')

# Subplot 2
plt.subplot(1, 2, 2)
sns.boxplot(data=df_mod, y='EmploymentStatus', x='BorrowerRate', palette='Greens_r')
plt.ylabel('Employment Status')
plt.xlabel('Borrower Rate')

plt.suptitle('Borrower Rate by Income Range and Employment Status')
plt.show()
```



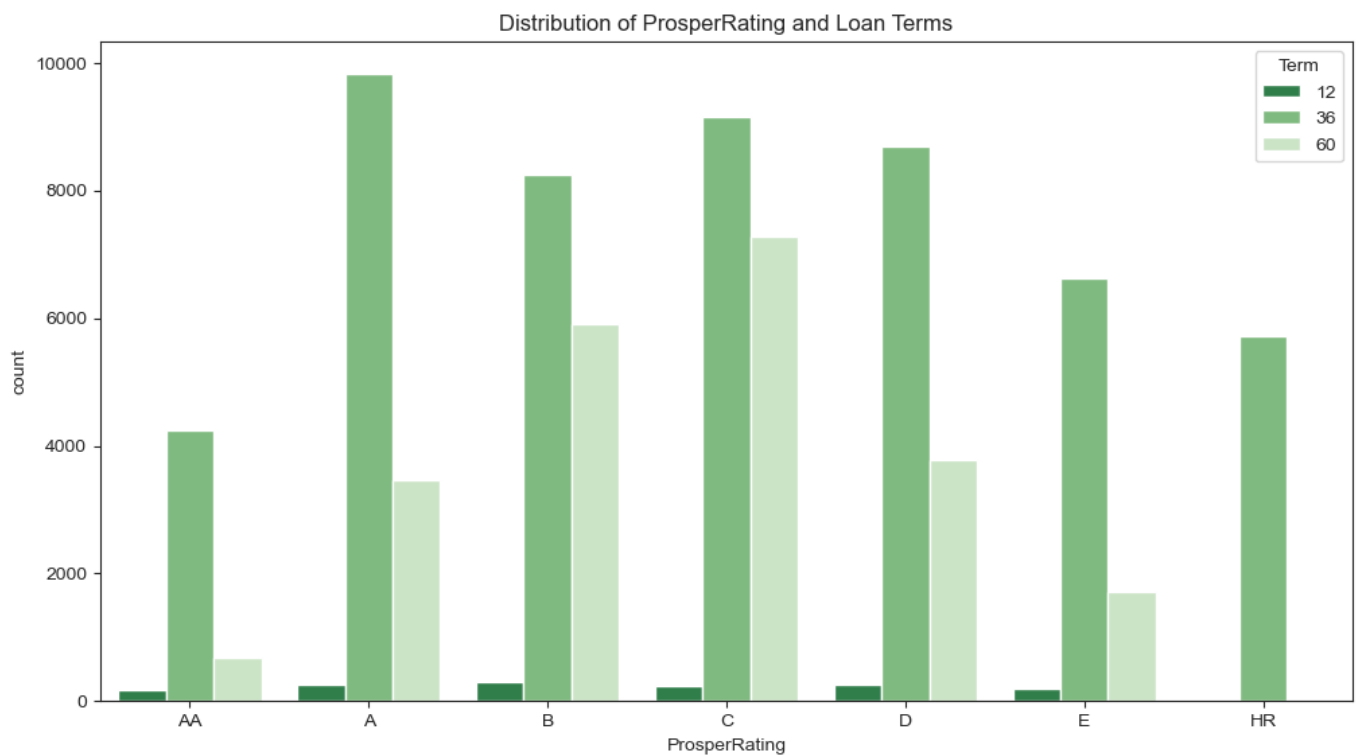
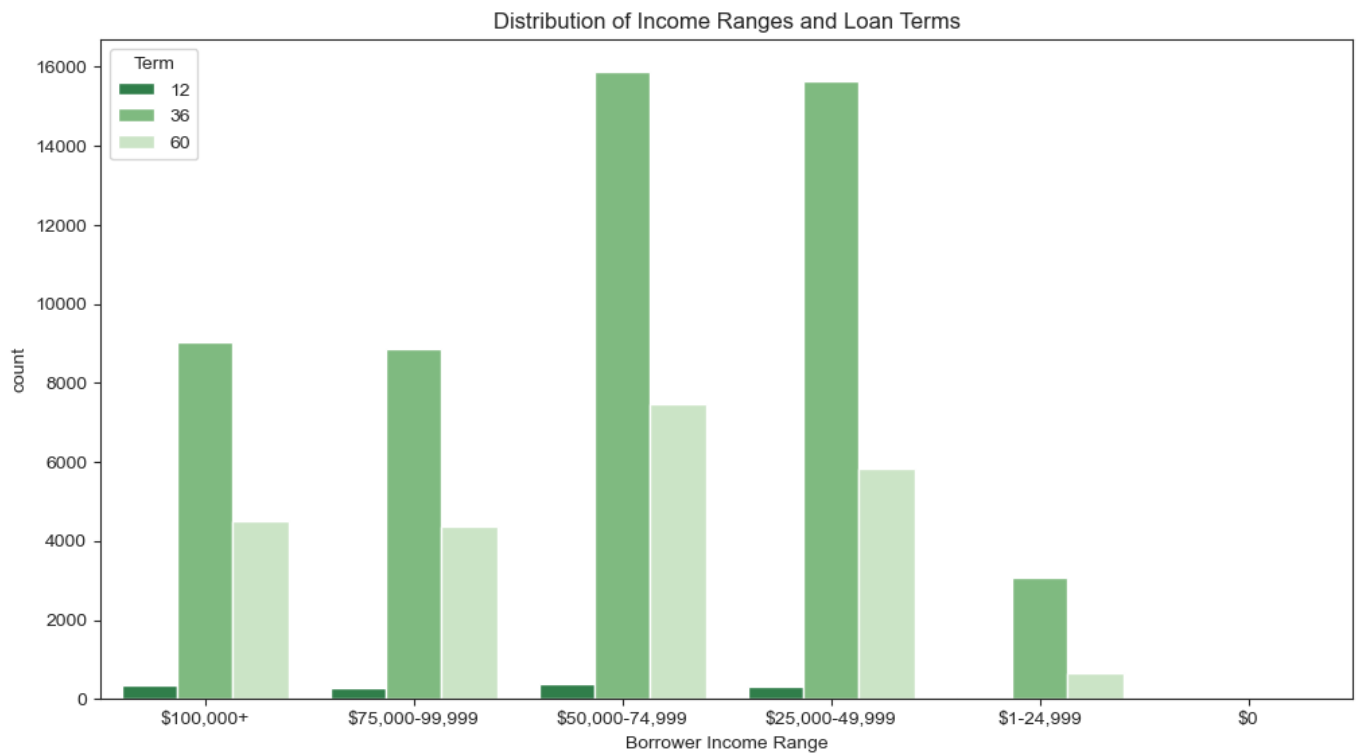
- The boxplots reveal that the median borrower rate decrease with higher income levels. The income range of 1-24,999k exhibits the highest median borrower rate.
- Within the employment status categories, the "unemployed" group displays the highest median borrower rate, indicating that unemployed individuals tend to incur higher interest costs on their loans.

3. Qualitative vs. Qualitative

```
In [57]: # Clustered bar chart of income range and loan term
plt.figure(figsize=[12, 14])

# Subplot 1: Income range and term distribution
plt.subplot(2, 1, 1)
sns.countplot(data=df_mod, x='IncomeRange', hue='Term', palette='Greens_r')
plt.title('Distribution of Income Ranges and Loan Terms')
plt.xlabel('Borrower Income Range')

# Subplot 2: Prosper rating and term distribution
plt.subplot(2, 1, 2)
sns.countplot(data=df_mod, x='ProsperRating', hue='Term', palette='Greens_r', order=or)
plt.title('Distribution of ProsperRating and Loan Terms')
plt.xlabel('ProsperRating');
```



- Borrowers with an income range of **\$50,000 - \$74,999** predominantly opt for 36-month loan terms, followed closely by those with income ranges of **\$25,000 - \$49,999**. Borrowers with no reported income primarily choose 36-month loan terms.
- Among Prosper rating categories, B and C ratings have a higher proportion of 60-month loans. In contrast, HR-rated borrowers exclusively opt for 36-month terms, while A-rated borrowers show a preference for 36-month loans as well.

Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset?

- Borrower APR** and **Borrower Rate** exhibit a robust positive correlation, with a coefficient of 0.99, which is expected since a higher APR typically leads to borrowers paying a greater amount of interest on their loans.

- **Borrower APR** and the **Loan Original Amount** demonstrate an inverse correlation with a coefficient of -0.32 . The accompanying scatter plot visually corroborates this negative association, illustrating that as the loan amount rises, the APR tends to decline.

Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

- As the **Loan Original Amount** rises, there is a concurrent growth in the **Monthly Loan Payment** (positive correlation).
- **Loan Original Amount** and **Stated Monthly Income** shows a slight positive correlation, as denoted by a correlation coefficient of 0.30.
- Borrowers within the income range of \$50,000 - \$74,999 primarily opt for 36-month loan terms, followed by those in the \$25,000 - \$49,999 income range. Borrowers with zero reported income exclusively select 36-month terms.
- Notably, individuals categorized as unemployed demonstrate the highest median borrower rate, signifying that they pay higher interest rates on loans.

Multivariate Exploration


```
In [58]: df_mod.BorrowerAPR.describe()
```

```
Out [58]: count      76768.000000
          mean         0.223978
          std         0.079291
          min         0.045830
          25%         0.162590
          50%         0.215660
          75%         0.287800
          max         0.423950
          Name: BorrowerAPR, dtype: float64
```

```
In [59]: # Create a scatter plot to show the distribution of LoanOriginalAmount by BorrowerAPR
plt.figure(figsize=(14, 8))
sns.scatterplot(data=df_mod, x='LoanOriginalAmount', y='BorrowerAPR', hue='ProsperRating',
               alpha=0.7, edgecolor='none')

plt.xlabel('LoanOriginalAmount')
plt.ylabel('BorrowerAPR')
plt.title('Scatterplot with Color Gradient (ProsperRating): Borrower APR vs. Loan Amount')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

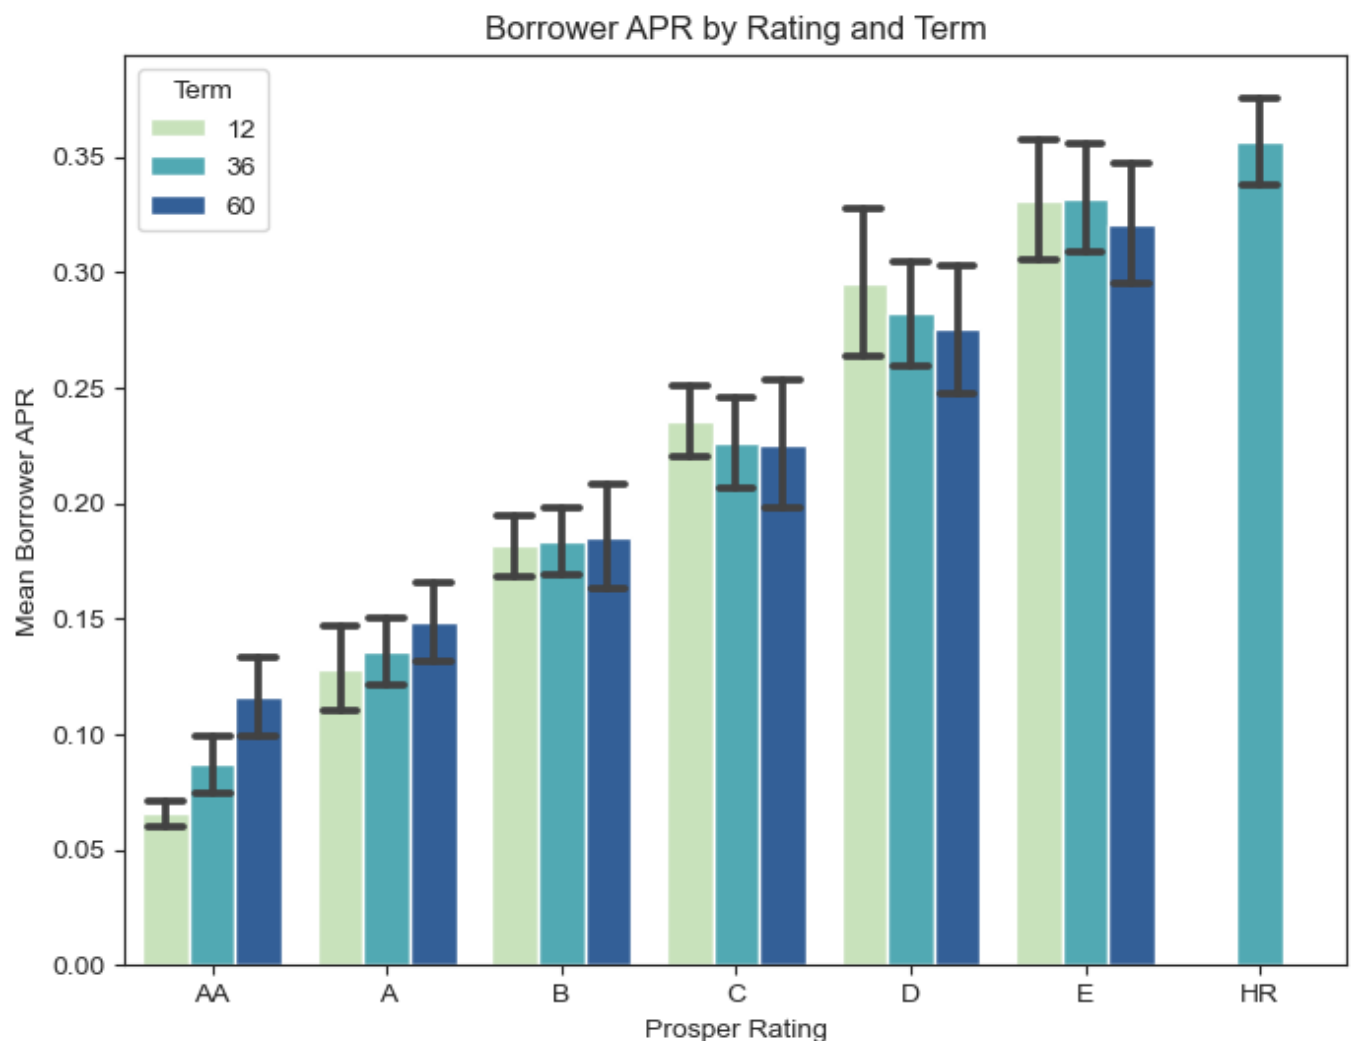


- As the **ProsperRating** improves, the **LoanOriginalAmount** tends to increase.
- With higher **ProsperRating**, borrowers generally experience lower **BorrowerAPR**.
- The correlation between **LoanOriginalAmount** and **BorrowerAPR** transitions from negative to slightly positive as **ProsperRating** improve.

```
In [60]: # Create a figure for the point plot
fig = plt.figure(figsize=[8, 6])

# Create a point plot to visualize Borrower APR across Prosper Rating and Loan Term
sns.barplot(data=df_mod, x='ProsperRating', y='BorrowerAPR', hue='Term', order=order,
            palette='YlGnBu', errorbar='sd', capsize=0.2)

# Set the title and labels for the plot
plt.title('Borrower APR by Rating and Term')
plt.xlabel('Prosper Rating')
plt.ylabel('Mean Borrower APR');
```



- **Higher ProsperRating (AA-B)** experience increasing APRs as the **Term** extends.
- Conversely, **lower ProsperRating (C-HR)** observe decreasing APRs as the **Term** extends.

Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest?

- As borrowers have better ratings, we observe that the loan amount tends to increase, while the borrower APR decreases. This suggests that borrowers with higher ratings receive larger loans at more favorable interest rates.
- When we specifically analyze the relationship between loan term and APR, interesting patterns emerge. Borrowers with excellent ratings (AA-B) experience increasing APRs as the loan term extends, which may seem counterintuitive. However, for lower-rated borrowers (C-HR), APRs tend to decrease as the loan term becomes longer. This suggests that loan terms interact differently with borrower ratings, potentially reflecting risk assessment variations by lenders.

Were there any interesting or surprising interactions between features?

- Notably, we found a negative correlation between BorrowerAPR and LoanOriginalAmount. This implies that as borrowers request larger loan amounts, the associated APR tends to decrease. This observation could be significant for both lenders and borrowers, as it indicates a potential benefit for borrowers seeking higher loan amounts in terms of lower interest rates.

- These relationships and interactions provide valuable insights into how borrower ratings, loan terms, loan amounts, and APRs are interconnected, allowing for a more comprehensive understanding of the lending dynamics in the dataset.

Conclusions

- **Borrower APR** and **Borrower Rate** exhibit a robust positive correlation, with a coefficient of 0.99, which is expected since a higher APR typically leads to borrowers paying a greater amount of interest on their loans.
- **Borrower APR** and the **Loan Original Amount** demonstrate an inverse correlation with a coefficient of -0.32 . The accompanying scatter plot visually corroborates this negative association, illustrating that as the loan amount rises, the APR tends to decline.
- **Higher ProsperRating (AA-B)** experience increasing **Borrower APR** as the **Term** extends.
- Conversely, **lower ProsperRating (C-HR)** observe decreasing **Borrower APR** as the **Term** extends.
- With higher **ProsperRating**, borrowers generally experience lower **BorrowerAPR**.
- The correlation between **LoanOriginalAmount** and **BorrowerAPR** transitions from negative to slightly positive as **ProsperRating** improve.

In []:

In []: